

1 What is the model (EGMN)?

The main proposed model in this repository is **EGMN** (Exponential–Gaussian Mixture Network), implemented in `model/egmn.py`. EGMN models watch time as a **mixture distribution** designed to capture both short “quick-skip” behavior and longer watch-time behavior.

1.1 Distribution family

EGMN uses a mixture with:

- **One Exponential component** (to capture the spike near zero watch time).
- **Multiple Gaussian components** (to capture longer watch time). In the code, the Gaussian components are implemented as **Normals truncated at 0** (to avoid negative time).

1.2 Architecture diagram




Figure 1: EGMN architecture overview: feature encoders + shared MLP + mixture-parameter heads, producing a distribution and a mean prediction.

If we denote the mixture weights by π , exponential rate by λ , and Gaussian parameters (μ_k, σ_k) , the model defines a mixture density over watch time $y \geq 0$ of the form:

$$p(y) = \pi_0 \text{Exp}(y; \lambda) + \sum_{k=1}^K \pi_k \text{TN}(y; \mu_k, \sigma_k, \text{lower} = 0),$$

where $\text{TN}(\cdot)$ is a Normal distribution truncated below at 0.

1.3 Neural architecture

- **Inputs:** a feature dictionary produced by the dataloader (sparse IDs, sequence features with masks, and continuous features such as `duration`).
- **Feature encoding:** embeddings for sparse/sequence features, and linear transforms for continuous features.
- **Shared trunk:** a multi-layer perceptron (MLP) over the concatenated feature representation.
- **Output heads:** predict mixture logits (converted to weights by softmax), the exponential rate λ (via softplus), and Gaussian parameters μ, σ (also via softplus for positivity constraints in the implementation).

1.4 Code snippet: heads and parameterization

```
# model/egmn.py (illustrative excerpt)
self.lambda_layer = nn.Sequential(
    nn.Linear(hidden_dim, 1),
    nn.Softplus(beta=0.5),
)

comp_num = 10 # number of Gaussian components
self.mixture_logits = nn.Linear(hidden_dim, comp_num + 1) # +1 for Exponential
self.gauss_mu = nn.Sequential(nn.Linear(hidden_dim, comp_num), nn.Softplus())
self.gauss_sigma = nn.Sequential(nn.Linear(hidden_dim, comp_num), nn.Softplus())
```

1.5 Training objective (as implemented)

Training uses a combination of:

- **Negative log-likelihood (NLL)** of the observed watch time under the mixture.
- **Reconstruction regularizer:** an L1 loss between the observed label y and the model’s mixture *mean* prediction.
- **Mixture entropy term:** computed from the mixture weights π , weighted by a hyperparameter (`--alpha` in `run_egmn.py`).

In the training script, the total loss is:

$$\mathcal{L} = \mathcal{L}_{\text{NLL}} + \alpha \mathcal{L}_{\text{entropy}} + \beta \mathcal{L}_{\text{reg}},$$

with α and β controlled by command-line flags.

1.6 Prediction used for evaluation

For evaluation, the code uses a **point prediction equal to the mixture mean**. In the implementation, the predicted value is computed as a weighted sum of component means (including the exponential mean $1/\lambda$).

1.7 Code snippet: prediction as mixture mean

```
# model/egmm.py (illustrative excerpt)
pi, lambda_, mu, sigma = self.forward(features)
pi = torch.softmax(pi, dim=1)

# component means: Exponential mean is 1/lambda; Gaussians use mu
component_means = torch.concat([1 / lambda_, mu], dim=1)
pred = torch.sum(pi * component_means, dim=1)
```